

# Novels by Dan Brown Textmining - Shiny

This R Markdown document is made interactive using Shiny.

## Make Packages Available

Here I load all the packages necessary for this project. The first thing I need to do is load some packages that I'm going to be using. I use pacman simply to manage packages, and tm is a text mining and that will give us most of our functionality. SnowballC adds some additional text analysis, and dplyr is for manipulating data and for arranging the code using pipes, where the output of one command feeds directly into the input of another one.

```
library(shiny)
library(wordcloud)
library(devtools)
library(tidyverse)
library(stringr)
library(tidytext)
library(dplyr)
library(reshape2)
library(igraph)
library(ggraph)
library(memoise)
if (packageVersion("devtools") < 1.6) {
  install.packages("devtools")
}
pacman::p_load(pacman, tm, SnowballC, dplyr)
```

## Import Three Books

The books this project is going to do textmining and sentiment analysis on are three novels by Dan Brown - "Angels & Demons", "The Da Vinci Code", and "The Lost Symbol". I'll start by importing book data, which is the full content of the three books. I have everything in the same directory, so there's no need to give a specific file path. I've already removed the metadata at the beginning and the end of the documents, so all that's left is the novels themselves.

```
# "Angels & Demons" by Dan Brown, published 2000
bookAAD <- readLines('ANGELS AND DEMONS.txt')

# "The Da Vinci Code" by Dan Brown, published 2003
bookDVC <- readLines('The Da Vinci Code.txt')

# "The Lost Symbol" by Dan Brown, published 2009
bookTLS <- readLines('The Lost Symbol.txt')
```

I'll begin by giving the data of every single book respectively, and then I'll compare their features in a set of 2 books and 3 later. First, I'm going to create a Corpus, which is a body of text for each book. I'll begin by creating what I call a preliminary corpus, because I'm going to do some later clean-up on it. These commands come from tm, for text mining. I'm going to remove the punctuation, any numbers, change everything to lowercase, and remove stopwords. Stopwords are words such as "the", "I", "but", which are usually meaningless when doing text mining.

I'm also going to stem the documents, and what that does is it takes a word like, "Stop" and it takes the variations of it, "Stops, stopped, stopping," and it cuts off those end parts and leaves us with just the beginning, "Stop."

```
# CORPUS FOR ANGELS & DEMONS

# Preliminary corpus
corpusAAD <- Corpus(VectorSource(bookAAD)) %>%
  tm_map(removePunctuation) %>%
  tm_map(removeNumbers) %>%
  tm_map(content_transformer(tolower)) %>%
  tm_map(removeWords, stopwords("english")) %>%
  tm_map(stripWhitespace) %>%
  tm_map(stemDocument)

# Create term-document matrices & remove sparse terms
tdmAAD <- DocumentTermMatrix(corpusAAD) %>%
  removeSparseTerms(1 - (5/length(corpusAAD)))
```

## Word Frequencies

Now I'm going to get absolute frequencies for each word, and then relative frequencies

```
# Calculate and sort by word frequencies
word.freqAAD <- sort(colSums(as.matrix(tdmAAD)),
                      decreasing = T)

# Create frequency table
tableAAD <- data.frame(word = names(word.freqAAD),
                        absolute.frequency = word.freqAAD,
                        relative.frequency =
                          word.freqAAD/length(word.freqAAD))

# Remove the words from the row names
rownames(tableAAD) <- NULL

# Show the 10 most common words
head(tableAAD, 10)
```

##	word	absolute.frequency	relative.frequency
## 1	angel	211	0.84
## 2	god	126	0.50
## 3	heaven	71	0.28
## 4	earth	62	0.25
## 5	bodi	55	0.22
## 6	will	51	0.20
## 7	one	42	0.17
## 8	bibl	42	0.17
## 9	know	40	0.16
## 10	time	39	0.15

As we can see in the table, Dan Brown uses "angel" 211 times and the relative frequency of "angel" is about 0.84.

I'm now going to create a csv file that has the most common words together with their absolute and relative frequencies. The file name will be AAD\_1000 in which AAD stands for Angels And Demons. The file will be saved to the same directory where I have my other documents.

```
# Export the 1000 most common words in CSV files
write.csv(tableAAD[1:1000, ], "AAD_1000.csv")
```

I'll repeat the same steps described above on the other two books in the following codes.

```
# CORPUS FOR THE DA VINCI CODE

corpusDVC <- Corpus(VectorSource(bookDVC)) %>%
  tm_map(removePunctuation) %>%
  tm_map(removeNumbers) %>%
  tm_map(content_transformer(tolower)) %>%
  tm_map(removeWords, stopwords("english")) %>%
  tm_map(stripWhitespace) %>%
  tm_map(stemDocument)
tdmDVC <- DocumentTermMatrix(corpusDVC) %>%
  removeSparseTerms(1 - (5/length(corpusDVC)))
word.freqDVC <- sort(colSums(as.matrix(tdmDVC)),
  decreasing = T)
tableDVC <- data.frame(word = names(word.freqDVC),
  absolute.frequency = word.freqDVC,
  relative.frequency =
    word.freqDVC/length(word.freqDVC))
rownames(tableDVC) <- NULL
head(tableDVC, 10)
```

##	word	absolute.frequency	relative.frequency
## 1	langdon	1579	0.61
## 2	sophi	1127	0.43
## 3	teab	601	0.23
## 4	said	536	0.21
## 5	now	430	0.17
## 6	look	418	0.16
## 7	fach	398	0.15
## 8	one	325	0.13
## 9	back	295	0.11
## 10	grail	290	0.11

```
write.csv(tableDVC[1:1000, ], "DVC_1000.csv")
```

```
# CORPUS FOR THE LOST SYMBOL
```

```
corpusTLS <- Corpus(VectorSource(bookTLS)) %>%
  tm_map(removePunctuation) %>%
  tm_map(removeNumbers) %>%
  tm_map(content_transformer(tolower)) %>%
  tm_map(removeWords, stopwords("english")) %>%
  tm_map(stripWhitespace) %>%
  tm_map(stemDocument)
tdmTLS <- DocumentTermMatrix(corpusTLS) %>%
  removeSparseTerms(1 - (5/length(corpusTLS)))
word.freqTLS <- sort(colSums(as.matrix(tdmTLS)),
  decreasing = T)
tableTLS <- data.frame(word = names(word.freqTLS),
  absolute.frequency = word.freqTLS,
  relative.frequency =
    word.freqTLS/length(word.freqTLS))
rownames(tableTLS) <- NULL
head(tableTLS, 10)
```

##	word	absolute.frequency	relative.frequency
## 1	langdon	1365	0.50
## 2	katherin	762	0.28
## 3	said	696	0.25
## 4	now	555	0.20
## 5	peter	553	0.20
## 6	man	441	0.16
## 7	look	437	0.16
## 8	pyramid	415	0.15
## 9	solomon	394	0.14
## 10	one	389	0.14

```
write.csv(tableTLS[1:1000, ], "TLS_1000.csv")
```

Here's the part where I'll compare their features in a set of 2 books to find out the most distinctive words. I'm going to create one called dProp, which is for a difference in proportions. Now, in this case, I'm simply taking the difference, a subtraction.

"Angels & Demons" vs "The Da Vinci Code"

```
# MOST DISTINCTIVE WORDS #####
```

```
# Set number of digits for output
```

```
options(digits = 2)
```

```
# Compare relative frequencies (via subtraction)
```

```
# ("Angels & Demons" vs "The Da Vinci Code")
```

```
AADvsDVC <- tableAAD %>%
```

```
  merge(tableDVC, by = "word") %>%
```

```
  mutate(dProp =
```

```
    relative.frequency.x -
```

```
    relative.frequency.y,
```

```
    dAbs = abs(dProp)) %>%
```

```
  arrange(desc(dAbs)) %>%
```

```
  rename(AAD.freq = absolute.frequency.x,
```

```
         AAD.prop = relative.frequency.x,
```

```
         DVC.freq = absolute.frequency.y,
```

```
         DVC.prop = relative.frequency.y)
```

```
# Show the 10 most distinctive terms
```

```
head(AADvsDVC, 10)
```

##	word	AAD.freq	AAD.prop	DVC.freq	DVC.prop	dProp	dAbs
## 1	angel	211	0.837	6	0.0023	0.83	0.83
## 2	god	126	0.500	105	0.0404	0.46	0.46
## 3	heaven	71	0.282	22	0.0085	0.27	0.27
## 4	earth	62	0.246	36	0.0139	0.23	0.23
## 5	bodi	55	0.218	76	0.0293	0.19	0.19
## 6	said	11	0.044	536	0.2065	-0.16	0.16
## 7	bibl	42	0.167	26	0.0100	0.16	0.16
## 8	creat	33	0.131	25	0.0096	0.12	0.12
## 9	will	51	0.202	228	0.0878	0.11	0.11
## 10	lord	28	0.111	18	0.0069	0.10	0.10

```
# Save full table to CSV
```

```
write.csv(AADvsDVC, "AAD vs DVC.csv")
```

As we can see in the table above, “angel” appears 211 times in Angels & Demons, while it only appears 6 times in The Da Vinci Code, which makes sense given the story. That’s why it has a positive dProp, or difference in proportions. The full table is going to be saved as a csv file in the same directory.

I’ll continue the same steps for the other two sets.

“Angels & Demons” vs “The Lost Symbol”

```
# ("Angels & Demons" vs "The Lost Symbol")
AADvsTLS <- tableAAD %>%
  merge(tableTLS, by = "word") %>%
  mutate(dProp =
    relative.frequency.x -
    relative.frequency.y,
    dAbs = abs(dProp)) %>%
  arrange(desc(dAbs)) %>%
  rename(AAD.freq = absolute.frequency.x,
    AAD.prop = relative.frequency.x,
    TLS.freq = absolute.frequency.y,
    TLS.prop = relative.frequency.y)
head(AADvsTLS, 10)
```

##	word	AAD.freq	AAD.prop	TLS.freq	TLS.prop	dProp	dAbs
## 1	angel	211	0.837	13	0.0048	0.83	0.83
## 2	god	126	0.500	156	0.0571	0.44	0.44
## 3	heaven	71	0.282	48	0.0176	0.26	0.26
## 4	earth	62	0.246	63	0.0231	0.22	0.22
## 5	said	11	0.044	696	0.2549	-0.21	0.21
## 6	peter	5	0.020	553	0.2026	-0.18	0.18
## 7	bodi	55	0.218	142	0.0520	0.17	0.17
## 8	bibl	42	0.167	46	0.0168	0.15	0.15
## 9	will	51	0.202	227	0.0832	0.12	0.12
## 10	jesus	31	0.123	24	0.0088	0.11	0.11

```
write.csv(AADvsTLS, "AAD vs TLS.csv")
```

### "The Da Vinci Code" vs "The Lost Symbol"

```
# ("The Da Vinci Code" vs "The Lost Symbol")
DVCvsTLS <- tableDVC %>%
  merge(tableTLS, by = "word") %>%
  mutate(dProp =
    relative.frequency.x -
    relative.frequency.y,
    dAbs = abs(dProp)) %>%
  arrange(desc(dAbs)) %>%
  rename(DVC.freq = absolute.frequency.x,
    DVC.prop = relative.frequency.x,
    TLS.freq = absolute.frequency.y,
    TLS.prop = relative.frequency.y)
head(DVCvsTLS, 10)
```

```
##      word DVC.freq DVC.prop TLS.freq TLS.freq dProp dAbs
## 1  peter      13   0.0050    553   0.2026 -0.198 0.198
## 2  solomon     17   0.0065    394   0.1443 -0.138 0.138
## 3  pyramid     40   0.0154    415   0.1520 -0.137 0.137
## 4   mason     16   0.0062    320   0.1172 -0.111 0.111
## 5 langdon    1579   0.6082   1365   0.5000  0.108 0.108
## 6  church     236   0.0909     11   0.0040  0.087 0.087
## 7 ancient      68   0.0262    260   0.0952 -0.069 0.069
## 8     man     254   0.0978    441   0.1615 -0.064 0.064
## 9 brother     14   0.0054    166   0.0608 -0.055 0.055
## 10 teacher    146   0.0562      9   0.0033  0.053 0.053
```

```
write.csv(DVCvsTLS, "DVC vs TLS.csv")
```

Here's the part where I'll compare their features in a set of 3 books

```
# Three BOOKS DATA
titles <- c("Angels & Demons", "The Da Vinci Code", "The Lost Symbol")

books <- list(bookAAD, bookDVC, bookTLS)

##Each book is an array in which each value in the array is a chapter
series <- tibble()
for(i in seq_along(titles)) {

  temp <- tibble(chapter = seq_along(books[[i]]),
                 text = books[[i]]) %>%
    unnest_tokens(word, text) %>%
    ##Here we tokenize each chapter into words
    mutate(book = titles[i]) %>%
    select(book, everything())

  series <- rbind(series, temp)
}
# set factor to keep books in order of publication
series$book <- factor(series$book, levels = rev(titles))
series
```

```
## # A tibble: 311,918 x 3
##   book          chapter word
##   <fct>          <int> <chr>
## 1 Angels & Demons      2 angels
## 2 Angels & Demons      3 and
## 3 Angels & Demons      3 demons
## 4 Angels & Demons      5 their
## 5 Angels & Demons      5 nature
## 6 Angels & Demons      5 origin
## 7 Angels & Demons      5 mtntsfry
## 8 Angels & Demons      6 ond
## 9 Angels & Demons      6 classification
## 10 Angels & Demons     10 four
## # ... with 311,908 more rows
```

We can get counts for each word using the count function.

```
series %>% count(word, sort = TRUE)
```

```
## # A tibble: 17,039 x 2
##   word      n
##   <chr> <int>
## 1 the    22244
## 2 of      7648
## 3 a       7134
## 4 to      7114
## 5 and     6528
## 6 in      4476
## 7 he      4066
## 8 was     4059
## 9 his     3494
## 10 had    3072
## # ... with 17,029 more rows
```

Many of the words in the top 10 most common words are stopwords. I'm going to remove the stopwords here and make a word cloud.

```
# Remove stopwords and make a word cloud
series$book <- factor(series$book, levels = rev(titles))
series %>%
  anti_join(stop_words) %>%
  count(word) %>%
  with(wordcloud(word, n, max.words = 100))
```

```
## Joining, by = "word"
```



A word cloud of names and terms, with 'katherine' being the largest word. Other prominent words include 'sophie', 'pyramid', 'teabing', 'peter', 'sato', 'sauliere', 'solomon', 'moment', 'church', 'grail', 'building', 'phone', 'god', 'voice', 'angels', 'darkness', 'hand', 'robert', 'time', 'fache', 'bellamy', 'peter', 'sophie', 'sophisticated', 'teabing', 'peter', 'sophie'.

```
(hp_nrc <- series %>%
  inner_join(get_sentiments("nrc")) %>%
  group_by(book, chapter, sentiment))
```

```
## # A tibble: 69,460 x 4
## # Groups:   book, chapter, sentiment [48,451]
##   book          chapter word      sentiment
##   <fct>          <int> <chr>    <chr>
## 1 Angels & Demons      10 radio    positive
## 2 Angels & Demons      32 throne   positive
## 3 Angels & Demons      32 throne   trust
## 4 Angels & Demons      33 elders    positive
## 5 Angels & Demons      33 elders    trust
## 6 Angels & Demons      38 subject  negative
## 7 Angels & Demons      41 spirit     positive
## 8 Angels & Demons      45 heavenly anticipation
## 9 Angels & Demons      45 heavenly joy
## 10 Angels & Demons     45 heavenly positive
## # ... with 69,450 more rows
```

9/21

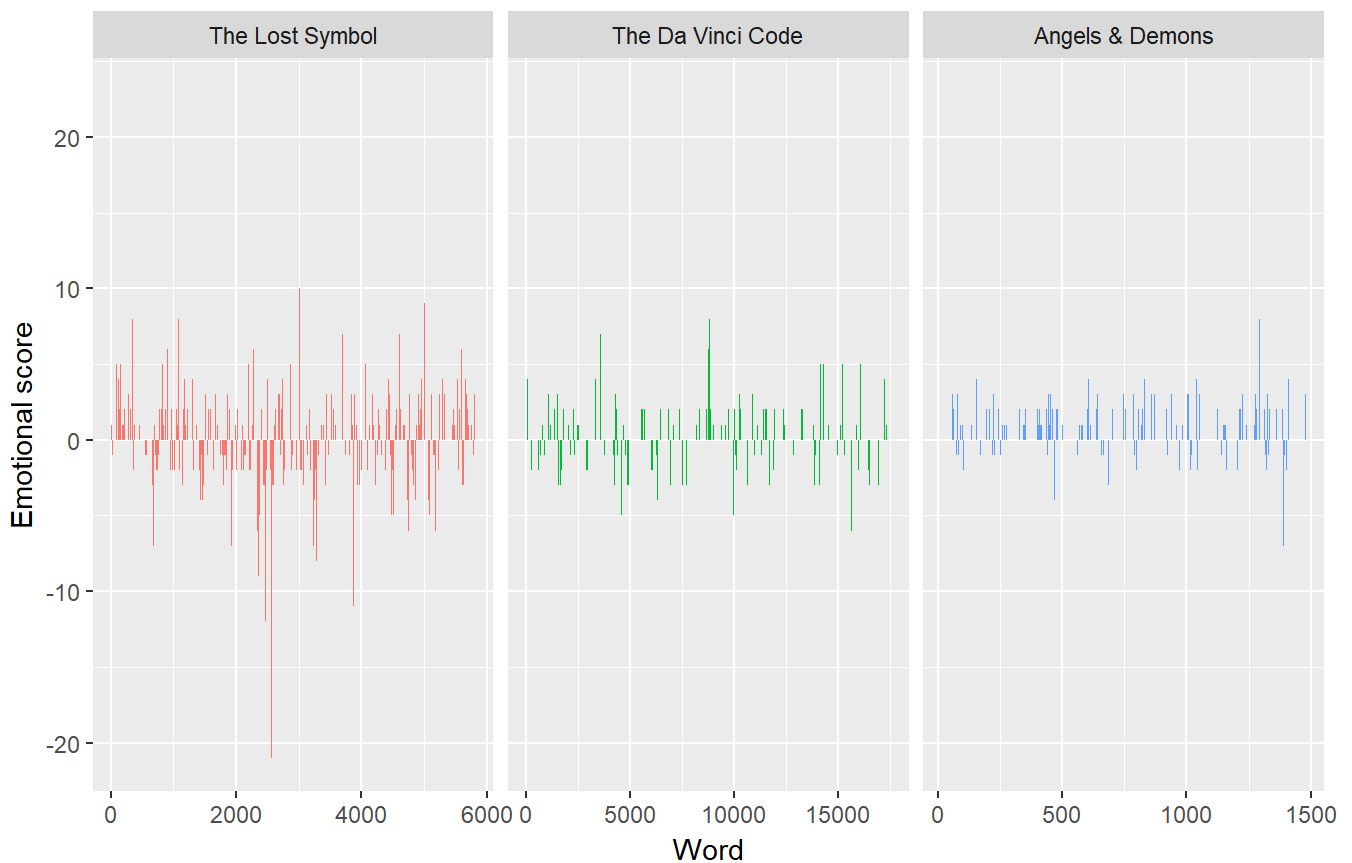
Here I want to visualize the positive/negative sentiment for each book over time using the AFINN dictionary.

```
series %>%
  inner_join(get_sentiments("afinn")) %>%
  group_by(book, chapter) %>%
  summarize(score = sum(score)) %>%
  ggplot(aes(chapter, score, fill = book)) +
  geom_col() +
  facet_wrap(~ book, scales = "free_x") +
  labs(title = "Emotional Arc of the Three Books",
       subtitle = "AFINN sentiment dictionary",
       x = "Word",
       y = "Emotional score") +
  theme(legend.position = "none")
```

```
## Joining, by = "word"
```

## Emotional Arc of the Three Books

AFINN sentiment dictionary




Here I'm going to show the cumulative score

```
# Cumulative score
series %>%
  inner_join(get_sentiments("afinn")) %>%
  group_by(book) %>%
  mutate(cumscore = cumsum(score)) %>%
  ggplot(aes(chapter, cumscore, fill = book)) +
  geom_step() +
  facet_wrap(~ book, scales = "free_x") +
  labs(title = "Emotional Arc of the Three Books",
       subtitle = "AFINN sentiment dictionary",
       x = "Word",
       y = "Cumulative emotional score")
```

```
## Joining, by = "word"
```

## Emotional Arc of the Three Books

AFINN sentiment dictionary



After this, I want to visualize the sentimental content of each chapter in each book using the NRC dictionary.


```

hp_nrc %>%
  count(sentiment, book, chapter) %>%
  filter(!(sentiment %in% c("positive", "negative"))) %>%
  # create area plot
  ggplot(aes(x = chapter, y = n)) +
  geom_col(aes(fill = sentiment)) +
  # add black smoothing line without standard error
  geom_smooth(aes(fill = sentiment), method = "loess", se = F, col = 'black') +
  theme(legend.position = 'none') +
  labs(x = "Word", y = "Emotion score", # add labels
       title = "Emotions during the books",
       subtitle = "Using tidytext and the nrc sentiment dictionary") +
  # seperate plots per sentiment and book and free up x-axes
  facet_grid(sentiment ~ book, scale = "free")

```

## Emotions during the books

Using tidytext and the nrc sentiment dictionary



### This chunk takes Longer to execute

## Another sentiment analysis

```

series %>%
  right_join(get_sentiments("nrc")) %>%
  filter(!is.na(sentiment)) %>%
  count(sentiment, sort = TRUE)

```

## Joining, by = "word"

```
## # A tibble: 10 x 2
##   sentiment      n
##   <chr>      <int>
## 1 positive   16491
## 2 negative   12039
## 3 trust      10829
## 4 anticipation 7330
## 5 fear       7103
## 6 joy        5749
## 7 sadness     5451
## 8 anger       4594
## 9 surprise    3763
## 10 disgust    3254
```

The 'bing' lexicon only classifies words as positive or negative.

```
series %>%
  right_join(get_sentiments("bing")) %>%
  filter(!is.na(sentiment)) %>%
  count(sentiment, sort = TRUE)
```

```
## Joining, by = "word"
```

```
## # A tibble: 2 x 2
##   sentiment      n
##   <chr>      <int>
## 1 negative   12204
## 2 positive    9978
```

Next I'm going to Use the the 'bing' lexicon for sentiment analysis and make a comparison cloud.

```
series %>%
  inner_join(get_sentiments("bing")) %>%
  count(word, sentiment, sort = TRUE) %>%
  acast(word ~ sentiment, value.var = "n", fill = 0) %>%
  comparison.cloud(colors = c("#F8766D", "#00BFC4"),
    max.words = 50)
```

```
## Joining, by = "word"
```

# negative



# positive

The comparison above still contains stopwords, and now I'm going to remove them and make a new cloud. I also use colors to separate words that are positive or negative. We can see here that character names don't appear in the following word cloud, because 'bing' doesn't classify names as positive or negative.

```
series %>%
  anti_join(stop_words) %>%
  inner_join(get_sentiments("bing")) %>%
  count(word, sentiment, sort = TRUE) %>%
  acast(word ~ sentiment, value.var = "n", fill = 0) %>%
  comparison.cloud(colors = c("#F8766D", "#00BFC4"),
    max.words = 50)
```

```
## Joining, by = "word"
## Joining, by = "word"
```

[illegible]

```
series %>%
  group_by(book) %>%
  mutate(word_count = 1:n(),
         index = word_count %% 500 + 1) %>%
  inner_join(get_sentiments("bing")) %>%
  count(book, index = index , sentiment) %>%
  ungroup() %>%
  spread(sentiment, n, fill = 0) %>%
  mutate(sentiment = positive - negative,
         book = factor(book, levels = titles)) %>%
  ggplot(aes(index, sentiment, fill = book)) +
  geom_bar(alpha = 0.5, stat = "identity", show.legend = FALSE) +
  facet_wrap(~ book, ncol = 2, scales = "free x")
```

15/21



Now, when we are doing sentiment analysis, using single words as tokens can be misleading sometimes. For example, if we say “This movie is not good,” we will know that this is a negative statement, while using single words as tokens can lead us to believe that this is a positive statement. To solve the problem, I’m going to consider pairs of words (bigrams). In the sentence “I love yellow cars,” the bigrams we can extract would be (I, love), (love, yellow), (yellow cars).

```
# Pairs of words (Bigrams)
series <- tibble()
for(i in seq_along(titles)) {

  temp <- tibble(chapter = seq_along(books[[i]]),
                 text = books[[i]]) %>%
    unnest_tokens(bigram, text, token = "ngrams", n = 2) %>%
    ##Here we tokenize each chapter into bigrams
    mutate(book = titles[i]) %>%
    select(book, everything())

  series <- rbind(series, temp)
}
# set factor to keep books in order of publication
series$book <- factor(series$book, levels = rev(titles))
series
```



```
## # A tibble: 301,131 x 3
##   book          chapter bigram
##   <fct>          <int> <chr>
## 1 Angels & Demons      1 <NA>
## 2 Angels & Demons      2 <NA>
## 3 Angels & Demons      3 and demons
## 4 Angels & Demons      4 <NA>
## 5 Angels & Demons      5 their nature
## 6 Angels & Demons      5 nature origin
## 7 Angels & Demons      5 origin mtntsfry
## 8 Angels & Demons      6 ond classification
## 9 Angels & Demons      7 <NA>
## 10 Angels & Demons     8 <NA>
## # ... with 301,121 more rows
```

I'm going to use the count function to find the most common bigrams in the books.

```
series %>%
  count(bigram, sort = TRUE)
```

```
## # A tibble: 133,203 x 2
##   bigram      n
##   <chr>   <int>
## 1 <NA>    7261
## 2 of the  2098
## 3 in the  1309
## 4 to the   988
## 5 on the   901
## 6 at the   691
## 7 had been 459
## 8 he had   440
## 9 into the 423
## 10 and the 416
## # ... with 133,193 more rows
```

As we can see there are still stopwords contained in the table, let's remove them.

```
# bigrams without stopwords
bigrams_separated <- series %>%
  separate(bigram, c("word1", "word2"), sep = " ")
bigrams_filtered <- bigrams_separated %>%
  filter(!word1 %in% stop_words$word) %>%
  filter(!word2 %in% stop_words$word)
# new bigram counts:
bigrams_united <- bigrams_filtered %>%
  unite(bigram, word1, word2, sep = " ")
bigrams_united %>%
  count(bigram, sort = TRUE)
```

```
## # A tibble: 35,369 x 2
##   bigram          n
##   <chr>          <int>
## 1 NA NA          7261
## 2 robert langdon    140
## 3 peter solomon    130
## 4 holy grail       119
## 5 da vinci         91
## 6 masonic pyramid   83
## 7 opus dei         75
## 8 ancient mysteries 71
## 9 katherine solomon 71
## 10 cell phone       70
## # ... with 35,359 more rows
```


Now, I'll use bigrams to practice tf-idf (term frequency -inverse document frequency). Tf-idf is an analysis that seeks to identify how common a word is in a particular text, given how often it occurs in a group of texts.

```
# Use bigrams to practice tf-idf
# (term frequency -inverse document frequency)
bigram_tf_idf <- bigrams_united %>%
  count(book, bigram) %>%
  bind_tf_idf(bigram, book, n) %>%
  arrange(desc(tf_idf))
bigram_tf_idf
```

```
## # A tibble: 36,511 x 6
##   book          bigram          n      tf   idf tf_idf
##   <fct>          <chr>    <int>  <dbl> <dbl>  <dbl>
## 1 Angels & Demons genesis 1      12 0.0100  1.10 0.0110
## 2 Angels & Demons lord jesus    11 0.00920 1.10 0.0101
## 3 Angels & Demons fallen angels   8 0.00669 1.10 0.00735
## 4 The Lost Symbol peter solomon 130 0.00509 1.10 0.00559
## 5 Angels & Demons 1 verse         5 0.00418 1.10 0.00459
## 6 Angels & Demons verse 2         5 0.00418 1.10 0.00459
## 7 Angels & Demons angels food     4 0.00334 1.10 0.00367
## 8 Angels & Demons isaiah 45       4 0.00334 1.10 0.00367
## 9 Angels & Demons thou hast       4 0.00334 1.10 0.00367
## 10 The Lost Symbol masonic pyramid 83 0.00325 1.10 0.00357
## # ... with 36,501 more rows
```

```
# Make the chart
plot <- bigram_tf_idf %>%
  arrange(desc(tf_idf)) %>%
  mutate(bigram = factor(bigram, levels = rev(unique(bigram))))
plot %>%
  top_n(20) %>%
  ggplot(aes(bigram, tf_idf, fill = book)) +
  geom_col() +
  labs(x = NULL, y = "tf-idf") +
  coord_flip()
```

```
## Selecting by tf_idf
```



I have now shown the bigrams, but in order to get an idea of how negations affect sentiment analysis, I want to find the bigrams that have the word “not” as the first word in the bigram to show how sentiment analysis was affected by negations,

```
bigrams_separated %>%
  filter(word1 == "not") %>%
  count(word1, word2, sort = TRUE)
```

```
## # A tibble: 445 x 3
##   word1 word2      n
##   <chr> <chr>   <int>
## 1 not    a         73
## 2 not    the        71
## 3 not    to         63
## 4 not    only        44
## 5 not    be         42
## 6 not    yet         41
## 7 not    even        39
## 8 not    sure        29
## 9 not    have        25
## 10 not   imagine     18
## # ... with 435 more rows
```

Removing stopwords,

```
# Remove stopwords
bigrams_separated <- bigrams_separated %>%
  filter(word1 == "not") %>%
  filter(!word2 %in% stop_words$word)%>%
  count(word1, word2, sort = TRUE)
bigrams_separated
```

```
## # A tibble: 289 x 3
##   word1 word2      n
##   <chr> <chr>   <int>
## 1 not   imagine    18
## 2 not   possibly    12
## 3 not   surprised    11
## 4 not   understand     9
## 5 not   coming        7
## 6 not   surprising     6
## 7 not   telling         6
## 8 not   begin          5
## 9 not   breathe         5
## 10 not   fathom          5
## # ... with 279 more rows
```

```
# Use Bing Lexicon
BING <- get_sentiments("bing")
not_words <- bigrams_separated %>%
  filter(word1 == "not") %>%
  filter(!word2 %in% stop_words$word)%>%
  inner_join(BING, by = c(word2 = "word")) %>%
  ungroup()
not_words
```


```
## # A tibble: 51 x 4
##   word1 word2      n sentiment
##   <chr> <chr>   <int> <chr>
## 1 not   entrust     3 positive
## 2 not   fail        3 negative
## 3 not   mistaken    3 negative
## 4 not   trust        3 positive
## 5 not   easy         2 positive
## 6 not   fall         2 negative
## 7 not   happy        2 positive
## 8 not   limited      2 negative
## 9 not   lying         2 negative
## 10 not   miss          2 negative
## # ... with 41 more rows
```


## Inputs and Outputs

Number of bins:

20

Bandwidth adjustment:





# Embedded Application

## Tabsets

**Distribution type:**

☒ Normal

☐ Uniform

☐ Log-normal

☐ Exponential

**Number of observations:**

1 500 1,000

1 101 201 301 401 501 601 701 801 901 1,000

Plot Summary Table

